

# Project1\_VenkataKanaparthi

January 9, 2022

## 0.1 Problem Statement

The terms “overweight” and “obesity” refer to body weight that is greater than what is considered normal or healthy for a certain height. Overweight is generally due to extra body fat. However, overweight may also be due to extra muscle, bone, or water. This would impact the health and there are many other diseases caused by extra body fat. We would be creating a model to predict the body fat based on some parameters. We would like to determine whether the person is obese or not in this project.

```
[1]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scikitplot as skplt
from imblearn.over_sampling import SMOTE
from sklearn.feature_selection import VarianceThreshold
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif
from sklearn.dummy import DummyClassifier
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.metrics import \
    classification_report, confusion_matrix, ConfusionMatrixDisplay, roc_auc_score, accuracy_score
from sklearn.metrics import auc, make_scorer, precision_recall_curve, log_loss
from sklearn.model_selection import cross_val_score
from numpy import mean, std
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.covariance import EllipticEnvelope
from sklearn.ensemble import IsolationForest
from sklearn.decomposition import PCA
from sklearn.cross_decomposition import PLSRegression
from sklearn.preprocessing import PowerTransformer, Normalizer
from sklearn.feature_selection import mutual_info_regression
from sklearn.inspection import permutation_importance
from sklearn.linear_model import Ridge, Lasso, ElasticNet, LinearRegression
from sklearn.model_selection import train_test_split, cross_val_score, \
    LeaveOneOut
```

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import make_pipeline
from sklearn.compose import TransformedTargetRegressor
from sklearn.metrics import r2_score, mean_squared_error, make_scorer
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import skew, kurtosis
from tqdm import tqdm

```

```

[2]: # Load data into a dataframe
bodyfat_df = pd.read_csv("bodyfat.csv")
bodyfat_df.head(10)

```

```

[2]:   Density  BodyFat  Age  Weight  Height  Neck  Chest  Abdomen  Hip  Thigh  \
0   1.0708    12.3   23   154.25   67.75   36.2   93.1    85.2   94.5   59.0
1   1.0853     6.1   22   173.25   72.25   38.5   93.6    83.0   98.7   58.7
2   1.0414    25.3   22   154.00   66.25   34.0   95.8    87.9   99.2   59.6
3   1.0751    10.4   26   184.75   72.25   37.4  101.8    86.4  101.2   60.1
4   1.0340    28.7   24   184.25   71.25   34.4   97.3   100.0  101.9   63.2
5   1.0502    20.9   24   210.25   74.75   39.0  104.5    94.4  107.8   66.0
6   1.0549    19.2   26   181.00   69.75   36.4  105.1    90.7  100.3   58.4
7   1.0704    12.4   25   176.00   72.50   37.8   99.6    88.5   97.1   60.0
8   1.0900     4.1   25   191.00   74.00   38.1  100.9    82.5   99.9   62.9
9   1.0722    11.7   23   198.25   73.50   42.1   99.6    88.6  104.1   63.1

      Knee  Ankle  Biceps  Forearm  Wrist
0   37.3   21.9   32.0    27.4   17.1
1   37.3   23.4   30.5    28.9   18.2
2   38.9   24.0   28.8    25.2   16.6
3   37.3   22.8   32.4    29.4   18.2
4   42.2   24.0   32.2    27.7   17.7
5   42.0   25.6   35.7    30.6   18.8
6   38.3   22.9   31.9    27.8   17.7
7   39.4   23.2   30.5    29.0   18.8
8   38.3   23.8   35.9    31.1   18.2
9   41.7   25.0   35.6    30.0   19.2

```

## EDA and Outlier Detection

```

[3]: # Check the dimension of the table
print("The dimension of the table is: ", bodyfat_df.shape)

# What type of variables are in the table
print("Describe Data")
print(bodyfat_df.describe())

```

The dimension of the table is: (252, 15)

Describe Data

```

count      Density      BodyFat      Age      Weight      Height      Neck  \
count    252.000000    252.000000    252.000000    252.000000    252.000000    252.000000

```

mean	1.055574	19.150794	44.884921	178.924405	70.148810	37.992063
std	0.019031	8.368740	12.602040	29.389160	3.662856	2.430913
min	0.995000	0.000000	22.000000	118.500000	29.500000	31.100000
25%	1.041400	12.475000	35.750000	159.000000	68.250000	36.400000
50%	1.054900	19.200000	43.000000	176.500000	70.000000	38.000000
75%	1.070400	25.300000	54.000000	197.000000	72.250000	39.425000
max	1.108900	47.500000	81.000000	363.150000	77.750000	51.200000

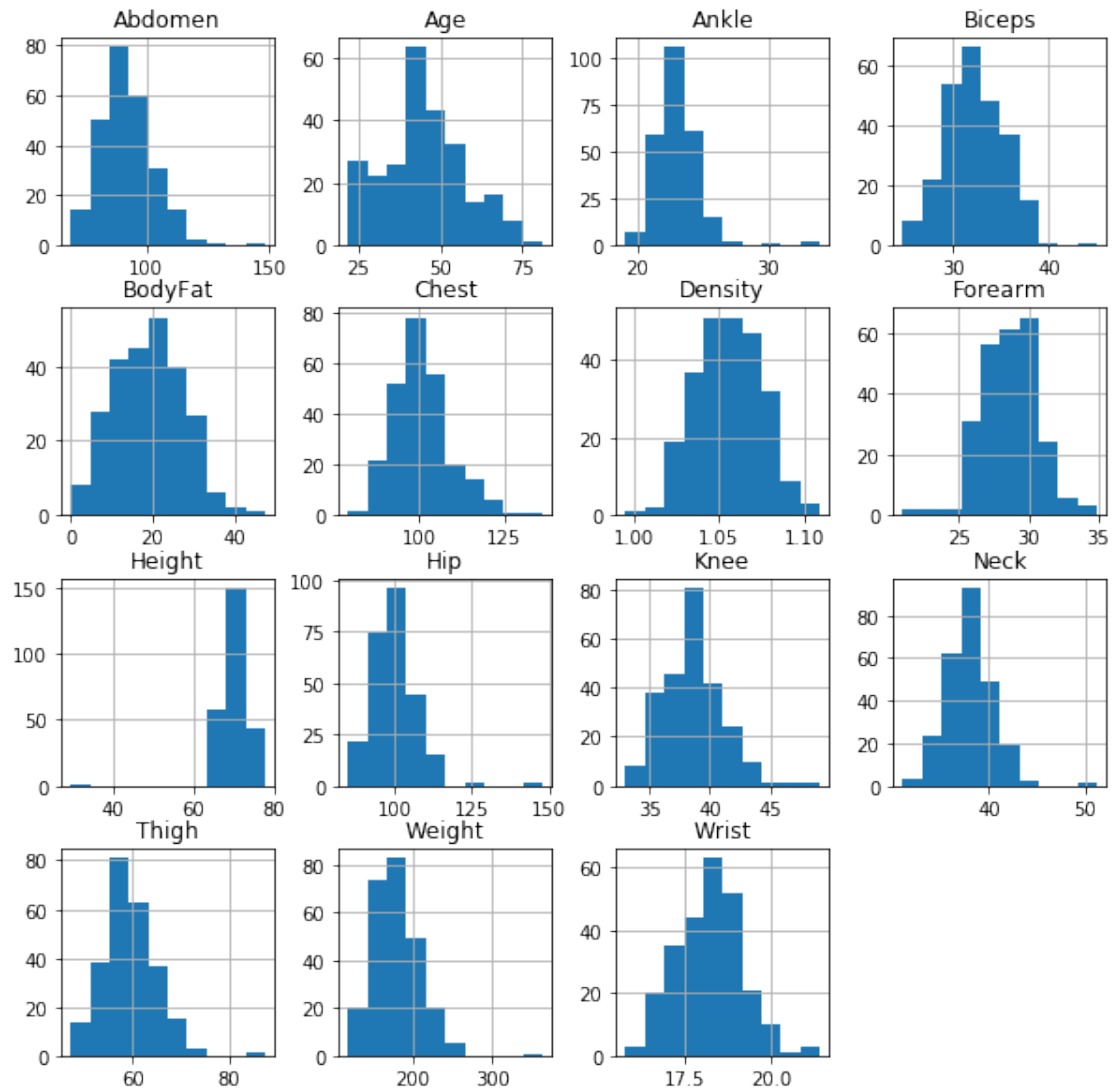
	Chest	Abdomen	Hip	Thigh	Knee	Ankle \
count	252.000000	252.000000	252.000000	252.000000	252.000000	252.000000
mean	100.824206	92.555952	99.904762	59.405952	38.590476	23.102381
std	8.430476	10.783077	7.164058	5.249952	2.411805	1.694893
min	79.300000	69.400000	85.000000	47.200000	33.000000	19.100000
25%	94.350000	84.575000	95.500000	56.000000	36.975000	22.000000
50%	99.650000	90.950000	99.300000	59.000000	38.500000	22.800000
75%	105.375000	99.325000	103.525000	62.350000	39.925000	24.000000
max	136.200000	148.100000	147.700000	87.300000	49.100000	33.900000

	Biceps	Forearm	Wrist
count	252.000000	252.000000	252.000000
mean	32.273413	28.663889	18.229762
std	3.021274	2.020691	0.933585
min	24.800000	21.000000	15.800000
25%	30.200000	27.300000	17.600000
50%	32.050000	28.700000	18.300000
75%	34.325000	30.000000	18.800000
max	45.000000	34.900000	21.400000


```
[4]: # Check if any missing values
      np.sum(np.sum(bodyfat_df.isna()))
```

```
[4]: 0
```

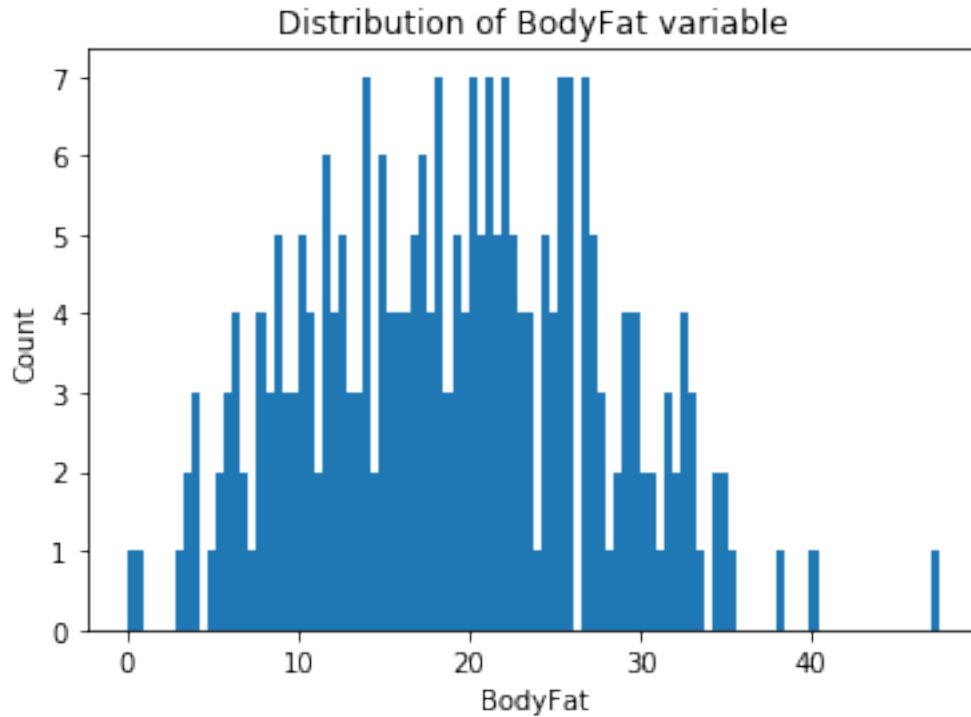
```
[5]: bodyfat_df.hist(figsize=(10,10))
      plt.show()
```



```
[6]: import warnings
warnings.filterwarnings('ignore')
plt.figure(figsize=(10,10))
sns.boxplot(data=bodyfat_df,color="white",linewidth=3)
sns.swarmplot(data=bodyfat_df,s=5,alpha=0.65)
plt.show()
```



```
[7]: # Plot histogram to identify any outlier that is visual to eye
plt.hist(bodyfat_df['BodyFat'], bins=100)
plt.ylabel('Count')
plt.xlabel('BodyFat')
plt.title('Distribution of BodyFat variable');
```



The values of 0 and 0.7 are not possible to achieve and we can delete them.

```
[8]: bodyfat_df.loc[bodyfat_df['BodyFat'] < 1]
```

```
[8]:
```


	Density	BodyFat	Age	Weight	Height	Neck	Chest	Abdomen	Hip	Thigh	\
171	1.0983	0.7	35	125.75	65.5	34.0	90.8	75.0	89.2	50.0	
181	1.1089	0.0	40	118.50	68.0	33.8	79.3	69.4	85.0	47.2	

	Knee	Ankle	Biceps	Forearm	Wrist
171	34.8	22.0	24.8	25.9	16.9
181	33.5	20.2	27.7	24.6	16.5

```
[9]: bodyfat_df1 = bodyfat_df.drop(bodyfat_df[bodyfat_df['BodyFat'] < 1].index)
```

```
[10]: # Plot histogram to identify any outlier that is visual to eye
plt.hist(bodyfat_df['Height'], bins=100)
plt.ylabel('Count')
plt.xlabel('Height')
plt.title('Distribution of Height variable');
```



```
[11]: bodyfat_df1.loc[bodyfat_df1['Height'] < 35]
```

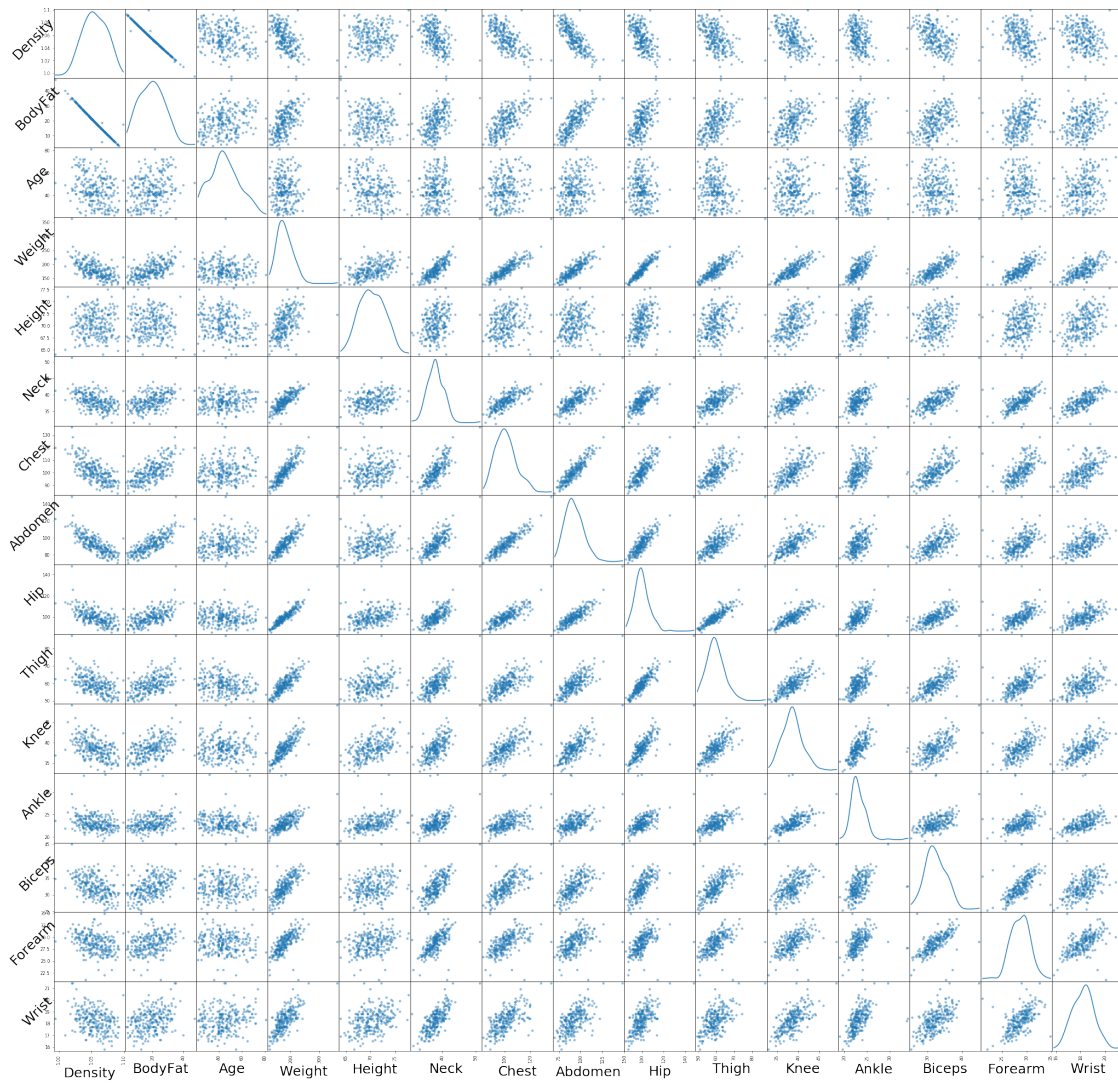
```
[11]:      Density  BodyFat  Age  Weight  Height  Neck  Chest  Abdomen  Hip  Thigh  \
41      1.025    32.9   44   205.0    29.5   36.6   106.0    104.3  115.5   70.6

      Knee  Ankle  Biceps  Forearm  Wrist
41  42.5   23.7   33.6    28.7   17.4
```

It looks like an outlier in the dataset, so I will delete this observation as well.

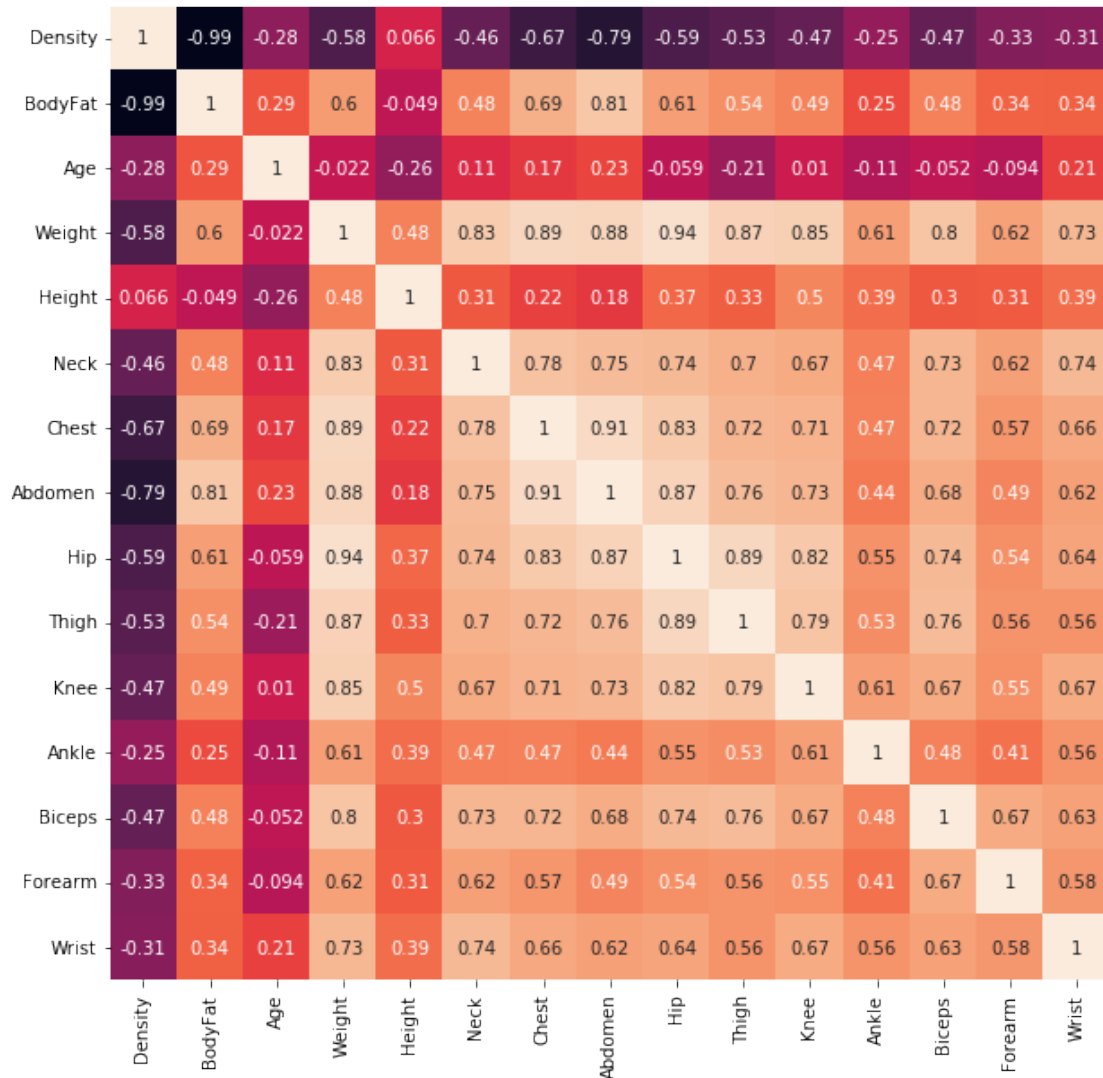
```
[12]: bodyfat_df1 = bodyfat_df1.drop(bodyfat_df1.loc[bodyfat_df1['Height'] < 35].
    ↪index)
```

```
[13]: # Scatter plot of the variables
axes = pd.plotting.scatter_matrix(bodyfat_df1, figsize=(35, 35), s=75,
    ↪diagonal='kde')
for ax in axes.flatten():
    ax.set_ylabel(ax.get_ylabel(), fontsize=25, rotation=45)
    ax.set_xlabel(ax.get_xlabel(), fontsize=25)
```



```
[14]: # Correlation of the variables
f,ax = plt.subplots(figsize=(10,10))
sns.heatmap(bodyfat_df1.corr(),annot=True,cbar=False,ax=ax)
plt.show()
```





Based on the scatter plot and the correlation heatmap it looks like there is a perfect linear relationship between Density and BodyFat variable. Since fat percentage and body density are synonymous, I will drop density and use only the circumference measurements to predict body fat percentage.

Also age doesn't determine the body fat of a person, so will drop Age variable also from the dataset

```
[15]: y = bodyfat_df1['BodyFat']
x = bodyfat_df1.drop(columns=['BodyFat', 'Density', 'Age'])
print(x)
```

	Weight	Height	Neck	Chest	Abdomen	Hip	Thigh	Knee	Ankle	Biceps	\
0	154.25	67.75	36.2	93.1	85.2	94.5	59.0	37.3	21.9	32.0	
1	173.25	72.25	38.5	93.6	83.0	98.7	58.7	37.3	23.4	30.5	
2	154.00	66.25	34.0	95.8	87.9	99.2	59.6	38.9	24.0	28.8	
3	184.75	72.25	37.4	101.8	86.4	101.2	60.1	37.3	22.8	32.4	

4	184.25	71.25	34.4	97.3	100.0	101.9	63.2	42.2	24.0	32.2
..	...	...	...	...	...	...	...	...	...	...
247	134.25	67.00	34.9	89.2	83.6	88.8	49.6	34.8	21.5	25.6
248	201.00	69.75	40.9	108.5	105.0	104.5	59.6	40.8	23.2	35.2
249	186.75	66.00	38.9	111.1	111.5	101.7	60.3	37.3	21.5	31.3
250	190.75	70.50	38.9	108.3	101.3	97.8	56.0	41.6	22.7	30.5
251	207.50	70.00	40.8	112.4	108.5	107.1	59.3	42.2	24.6	33.7

	Forearm	Wrist
0	27.4	17.1
1	28.9	18.2
2	25.2	16.6
3	29.4	18.2
4	27.7	17.7
..	...	...
247	25.7	18.5
248	28.6	20.1
249	27.2	18.0
250	29.4	19.8
251	30.0	20.9

[249 rows x 12 columns]

```
[16]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.
      ↪2,random_state=42)
```

```
[17]: # Details of training dataset
print("Shape of x_train dataset: ", x_train.shape)
print("Shape of y_train dataset: ", y_train.shape)
print("Shape of x_test dataset: ", x_test.shape)
print("Shape of y_test dataset: ", y_test.shape)
```

```
Shape of x_train dataset: (199, 12)
Shape of y_train dataset: (199,)
Shape of x_test dataset: (50, 12)
Shape of y_test dataset: (50,)
```

```
[18]: # Feature selection using SelectKBest feature selection
skbest = SelectKBest(k=10)
skbest.fit(x_train,y_train)
x_train_skbest=skbest.transform(x_train)
x_test_skbest=skbest.transform(x_test)
x_train_skbest.shape
```

```
[18]: (199, 10)
```

```
[19]: # 10 best features using SelectKBest
best_features = SelectKBest(score_func=f_classif, k=10)
```

```

fit = best_features.fit(x_train,y_train)
df_scores = pd.DataFrame(fit.scores_)
df_columns = pd.DataFrame(x_train.columns)
feature_scores = pd.concat([df_columns, df_scores],axis=1)
feature_scores.columns = ['Feature_Name','Score'] # name output columns
print(feature_scores.nlargest(10,'Score')) # print 10 best features

```

	Feature_Name	Score
4	Abdomen	3.809316
3	Chest	2.662361
5	Hip	2.475324
0	Weight	2.117462
2	Neck	1.948599
6	Thigh	1.768280
9	Biceps	1.545362
7	Knee	1.457433
1	Height	1.193745
11	Wrist	1.139133

Looks like Abdomen, Chest, Hip and Weight plays a major part when compared to the other features

## Model Evaluation

```

[20]: def get_model_score(pipe, data_x, data_y, cv=5):
        cv_score = cross_val_score(pipe, data_x, data_y, cv=cv)
        return np.mean(cv_score), np.std(cv_score)

def get_pipeline(*additional_pipe_steps, model):
    pipe_steps = [PowerTransformer(method='box-cox'), Normalizer(norm='l2')]
    pipe_steps.extend(additional_pipe_steps)
    return make_pipeline(*pipe_steps,
                          TransformedTargetRegressor(regressor=model,
                                                         ↵
                                                         ↪transformer=PowerTransformer(method='box-cox'))))

[21]: def evaluate_model(pipe, X, y):
        y_pred, y_true = np.empty(len(y)), np.empty(len(y))
        loo = LeaveOneOut()
        for i, (train_idx, test_idx) in tqdm(enumerate(loo.split(X)), total=len(y)):
            X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
            y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]
            y_pred[i] = pipe.fit(X_train, y_train).predict(X_test)[0]
            y_true[i] = y_test
        return r2_score(y_true, y_pred), np.sqrt(mean_squared_error(y_true, y_pred))

[22]: pipe = get_pipeline(PCA(n_components=6), model=LinearRegression())
        model_score, score_std = get_model_score(pipe=pipe, data_x=x, data_y=y, cv=5)
        r2, rmse = evaluate_model(pipe, x, y)

```

```
print('Linear Regression Model Score: ', model_score)
print('Linear Regression SD Score: ', score_std)
print('R Square: ', r2)
print('Root Mean Sqaure: ', rmse)
```

100%| | 249/249 [00:10<00:00, 24.40it/s]

Linear Regression Model Score: 0.4760618237463552

Linear Regression SD Score: 0.2588001654989154

R Square: 0.5842677076011253

Root Mean Sqaure: 5.277296230935636

```
[23]: pipe = get_pipeline(PCA(n_components=6), model=RandomForestRegressor())
model_score, score_std = get_model_score(pipe=pipe, data_x=x, data_y=y, cv=5)
r2, rmse = evaluate_model(pipe, x, y)
print('Random Forest Regressor Model Score: ', model_score)
print('Random Forest Regressor SD Score: ', score_std)
print('R Square: ', r2)
print('Root Mean Sqaure: ', rmse)
```

100%| | 249/249 [01:13<00:00, 3.40it/s]

Random Forest Regressor Model Score: 0.4713688065398777

Random Forest Regressor SD Score: 0.22240534030226908

R Square: 0.5808876913007635

Root Mean Sqaure: 5.298705727658999

Linear Regression is providing a better accuracy when compared with Random Forest. So, I would like to choose Linear Regression model in calculating the BodyFat for a provided data

```
[24]: def predict_loo(model, X, y):
    y_pred, y_true = np.empty(len(y)), np.empty(len(y))
    loo = LeaveOneOut()
    for i, (train_idx, test_idx) in tqdm(enumerate(loo.split(X)), total=len(y)):
        X_train, X_test = X.iloc[train_idx], X.iloc[test_idx]
        y_train, y_test = y.iloc[train_idx], y.iloc[test_idx]
        y_pred[i] = model.fit(X_train, y_train).predict(X_test)[0]
        y_true[i] = y_test
    return y_pred, y_true
```

```
[25]: final_model = get_pipeline(model=LinearRegression())
y_pred, y_true = predict_loo(final_model, x, y)
```

100%| | 249/249 [00:09<00:00, 26.54it/s]

These are some random measurements which would be given as input to provide the bodyfat of a person

```
[26]: measurements = np.array([165, 180, 38.5, 102.5, 90, 55.2, 38.3, 26.3, 32.5, 29.  
    ↪4, 17.2, 12.3]).reshape(1, -1)
```

```
[27]: prediction = final_model.predict(measurements)[0]  
    print(f'The prediction for body fat percent is {prediction:.1f}%')
```

The prediction for body fat percent is 17.3%.