

Term Project_draft_2_v1

November 20, 2021

```
[1]: # Course      : DSC630 - Predictive Analytics
     # Project    : Credit Card Fraud Detection
     # Name       : Vasanthakumar Kalaikkovan, Ganeshkumar Muthusamy and Venkata_
     ↪ Kanaparthi
```

0.1 Problem Statement

Credit card fraud is a major problem in financial services and costs billions of dollars every year. Credit card fraud continues to increase due to the rise and acceleration of Phone Order / Mail Order / E-Commerce. There has been tremendous use of credit cards for online shopping which led to a high amount of fraud related to credit cards. Financial institutions like Visa, MasterCard, Amex, and all debit networks have mandated that banks and merchants introduce EMV card technology to counter the fraud. In 2018, a total of \$24.26 Billion was lost due to payment card fraud across the globe, and the USA is the most fraud-prone country. Credit card fraud was ranked the number one type of identity theft fraud. Credit card fraud increased by 18.4% in 2018 and is still climbing. There can be two kinds of card fraud, card-present fraud, and card-not-present fraud.

```
[2]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scikitplot as skplt
from sklearn.model_selection import train_test_split, GridSearchCV
from imblearn.over_sampling import SMOTE
from sklearn.feature_selection import VarianceThreshold
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif
from sklearn.dummy import DummyClassifier
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.metrics import_
    ↪ classification_report, confusion_matrix, ConfusionMatrixDisplay, roc_auc_score
from sklearn.metrics import auc, make_scorer, precision_recall_curve, log_loss
from sklearn.model_selection import cross_val_score
from numpy import mean, std
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn.linear_model import SGDClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
import scikitplot as skplt

```

```

[3]: # Load data into a dataframe
df = pd.read_csv("creditcard.csv")
df.head(10)

```

```

[3]:   Time      V1      V2      V3      V4      V5      V6      V7 \
0    0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1    0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2    1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3    1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4    2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941
5    2.0 -0.425966  0.960523  1.141109 -0.168252  0.420987 -0.029728  0.476201
6    4.0  1.229658  0.141004  0.045371  1.202613  0.191881  0.272708 -0.005159
7    7.0 -0.644269  1.417964  1.074380 -0.492199  0.948934  0.428118  1.120631
8    7.0 -0.894286  0.286157 -0.113192 -0.271526  2.669599  3.721818  0.370145
9    9.0 -0.338262  1.119593  1.044367 -0.222187  0.499361 -0.246761  0.651583

      V8      V9  ...      V21      V22      V23      V24      V25 \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -0.206010
5  0.260314 -0.568671  ... -0.208254 -0.559825 -0.026398 -0.371427 -0.232794
6  0.081213  0.464960  ... -0.167716 -0.270710 -0.154104 -0.780055  0.750137
7 -3.807864  0.615375  ...  1.943465 -1.015455  0.057504 -0.649709 -0.415267
8  0.851084 -0.392048  ... -0.073425 -0.268092 -0.204233  1.011592  0.373205
9  0.069539 -0.736727  ... -0.246914 -0.633753 -0.120794 -0.385050 -0.069733

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62      0
1  0.125895 -0.008983  0.014724   2.69      0
2 -0.139097 -0.055353 -0.059752  378.66      0
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0
5  0.105915  0.253844  0.081080    3.67      0
6 -0.257237  0.034507  0.005168    4.99      0
7 -0.051634 -1.206921 -1.085339   40.80      0
8 -0.384157  0.011747  0.142404   93.20      0
9  0.094199  0.246219  0.083076    3.68      0

```

[10 rows x 31 columns]

```
[4]: # Check the dimension of the table
print("The dimension of the table is: ", df.shape)
```

The dimension of the table is: (284807, 31)

```
[5]: # What type of variables are in the table
print("Describe Data")
print(df.describe())
```

Describe Data

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.919560e-15	5.688174e-16	-8.769071e-15	2.782312e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	-1.552563e-15	2.010663e-15	-1.694249e-15	-1.927028e-16	-3.137024e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	V21	V22	V23	V24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.537294e-16	7.959909e-16	5.367590e-16	4.458112e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	1.453003e-15	1.699104e-15	-3.660161e-16	-1.206049e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000

```
max      7.519589e+00  3.517346e+00  3.161220e+01  3.384781e+01  25691.160000
```

```
          Class
count  284807.000000
mean      0.001727
std       0.041527
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max       1.000000
```

```
[8 rows x 31 columns]
```

```
[6]: # Check if any missing values
      df.isnull().sum()
```

```
[6]: Time      0
      V1        0
      V2        0
      V3        0
      V4        0
      V5        0
      V6        0
      V7        0
      V8        0
      V9        0
      V10       0
      V11       0
      V12       0
      V13       0
      V14       0
      V15       0
      V16       0
      V17       0
      V18       0
      V19       0
      V20       0
      V21       0
      V22       0
      V23       0
      V24       0
      V25       0
      V26       0
      V27       0
      V28       0
      Amount    0
```

```
Class      0
dtype: int64
```

```
[7]: # Check the types of each feature
df.dtypes
```


```
[7]: Time      float64
V1          float64
V2          float64
V3          float64
V4          float64
V5          float64
V6          float64
V7          float64
V8          float64
V9          float64
V10         float64
V11         float64
V12         float64
V13         float64
V14         float64
V15         float64
V16         float64
V17         float64
V18         float64
V19         float64
V20         float64
V21         float64
V22         float64
V23         float64
V24         float64
V25         float64
V26         float64
V27         float64
V28         float64
Amount      float64
Class       int64
dtype: object
```

```
[8]: # Histograms fraudulent and non-fraudulent transactions
fraud = df[df.Class == 1]
non_fraud = df[df.Class == 0]
plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
fraud.Time.hist(color='#F08080', bins=50, label="Fraudulent transaction")
plt.legend()
```

```
plt.subplot(2, 2, 2)
non_fraud.Time.hist(color='#66CDAA', bins=50, label="Non-fraudulent_
↳transaction")
plt.legend()
```

[8]: <matplotlib.legend.Legend at 0x7f8929110b20>




```
[9]: # Pie Chart fraudulent % vs non-fraudulent %
my_data = "Class"
grouped = df[my_data].value_counts().reset_index()
grouped = grouped.rename(columns = {my_data : "count", "index" : my_data})

labels = grouped[my_data]
sizes = grouped['count']

fig, ax = plt.subplots()
ax.pie(sizes, labels=labels, autopct='%1.2f%%', shadow=True, startangle=0,
↳colors = ['#66CDAA', '#F08080'])
ax.axis('equal')
plt.title('0: Not Fraud, 1: Fraud', fontsize=14)
plt.show()
```

0: Not Fraud, 1: Fraud



```
[10]: # Density plot of the features
var = df.columns.values

x = 0
non_fraud = df.loc[df['Class'] == 0]
fraud = df.loc[df['Class'] == 1]

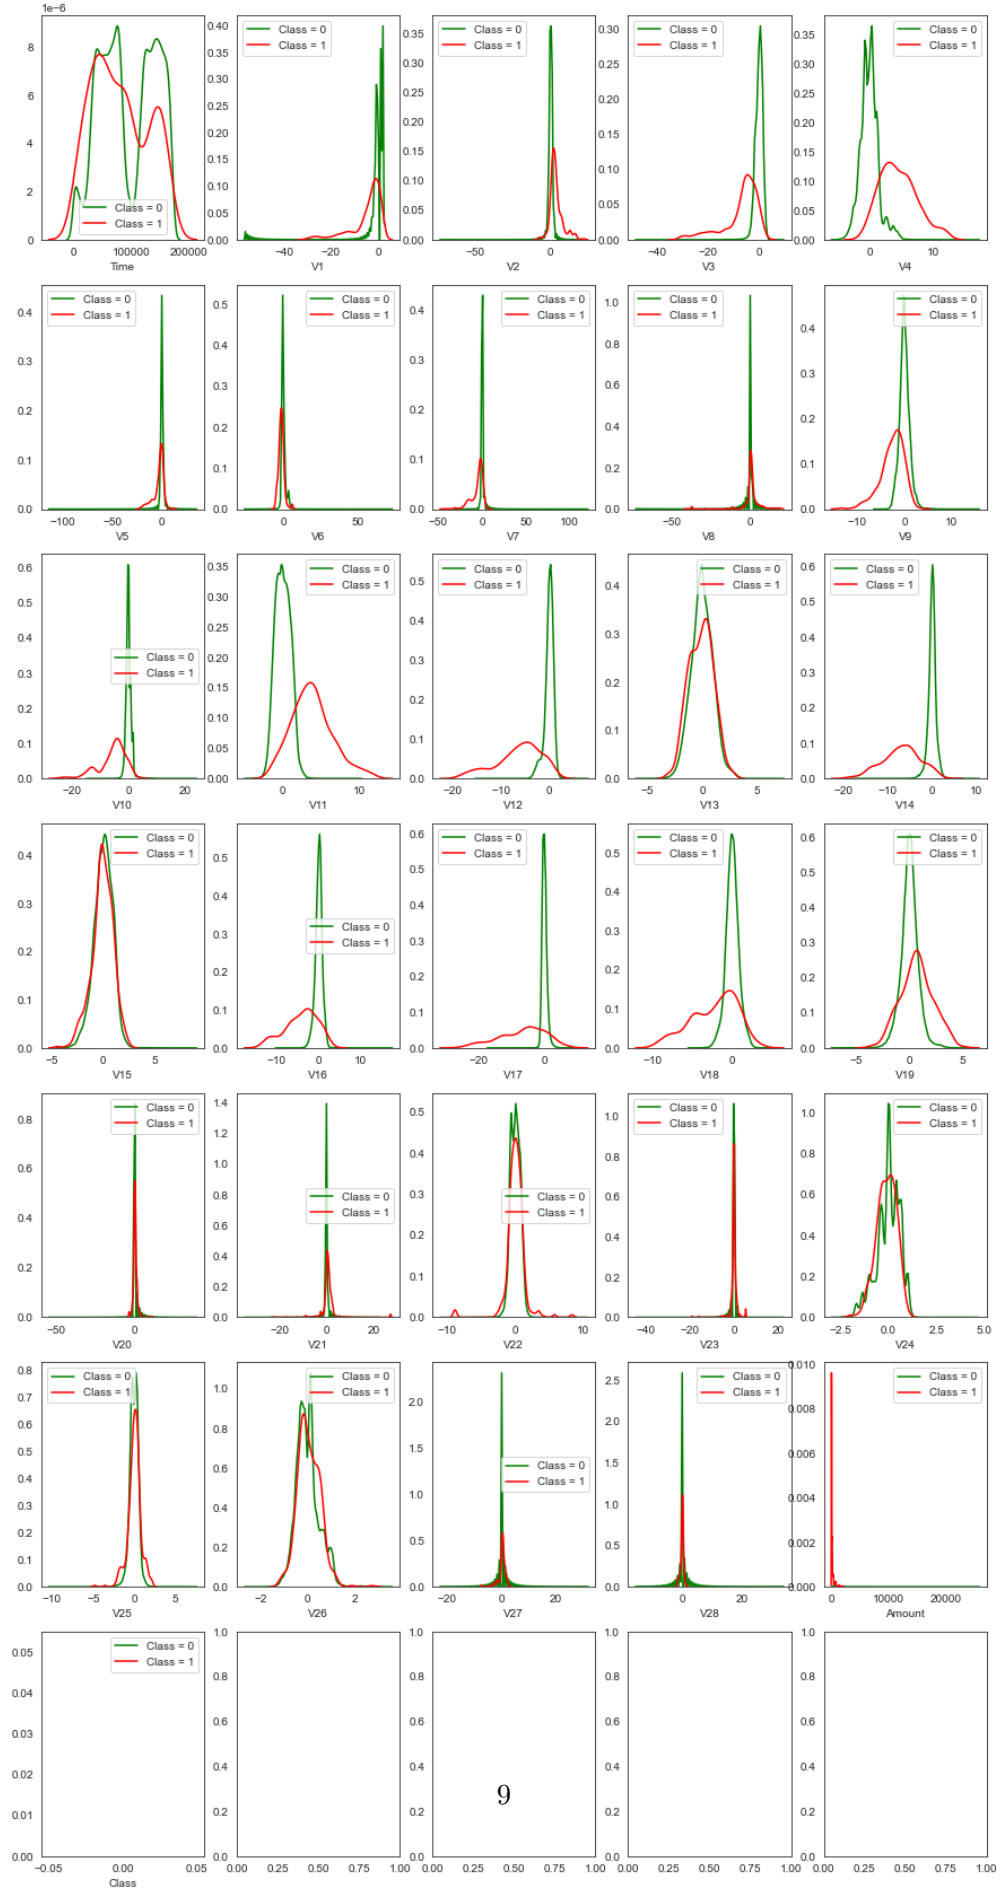
sns.set_style('white')
plt.figure()
fig, ax = plt.subplots(7,5,figsize=(15,30))

for feature in var:
    x += 1
    plt.subplot(7,5,x)
    sns.kdeplot(non_fraud[feature],label="Class = 0", color='green')
    sns.kdeplot(fraud[feature],label="Class = 1", color='red')
    plt.xlabel(feature, fontsize=10)
    locs, labels = plt.xticks()
    plt.tick_params(axis='both', which='major', labelsize=10)
plt.show();
```

```
/Users/ganeshkumar/opt/anaconda3/lib/python3.8/site-
packages/seaborn/distributions.py:283: UserWarning: Data must have variance to
compute a kernel density estimate.
```

```
warnings.warn(msg, UserWarning)
/Users/ganeshkumar/opt/anaconda3/lib/python3.8/site-
```

```
packages/seaborn/distributions.py:283: UserWarning: Data must have variance to
compute a kernel density estimate.
  warnings.warn(msg, UserWarning)
<Figure size 432x288 with 0 Axes>
```

```
[11]: def topn(df,n):
    npa = df.values

    npa = np.tril(npa, -1)
    topn_ind = np.argsortort(npa,-n,None)[-n:] #flatend ind, unsorted
    topn_ind = topn_ind[np.argsort(npa.flat[topn_ind])][::-1] #arg sort in
    ↪descending order
    cols,indx = np.unravel_index(topn_ind,npa.shape,'F') #unflatten, using
    ↪column-major ordering

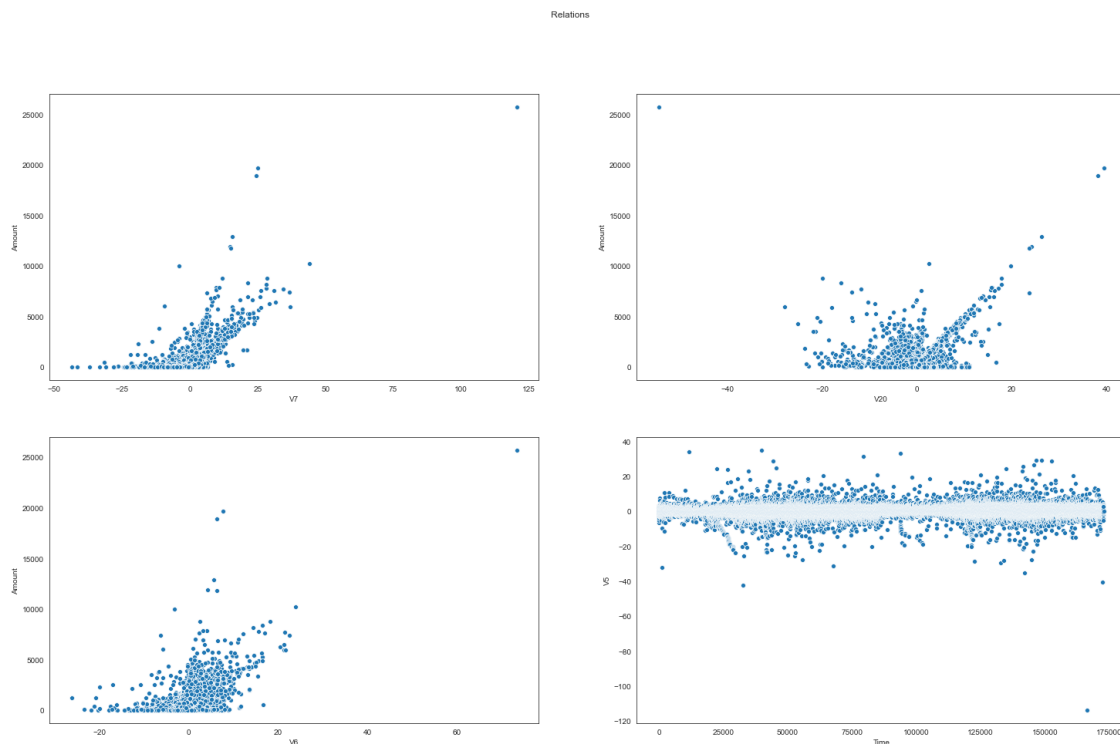
    return ([df.columns[c] for c in cols],[df.index[i] for i in indx])
```

```
[12]: max_corr_x , max_corr_y = topn(df.corr(),4)

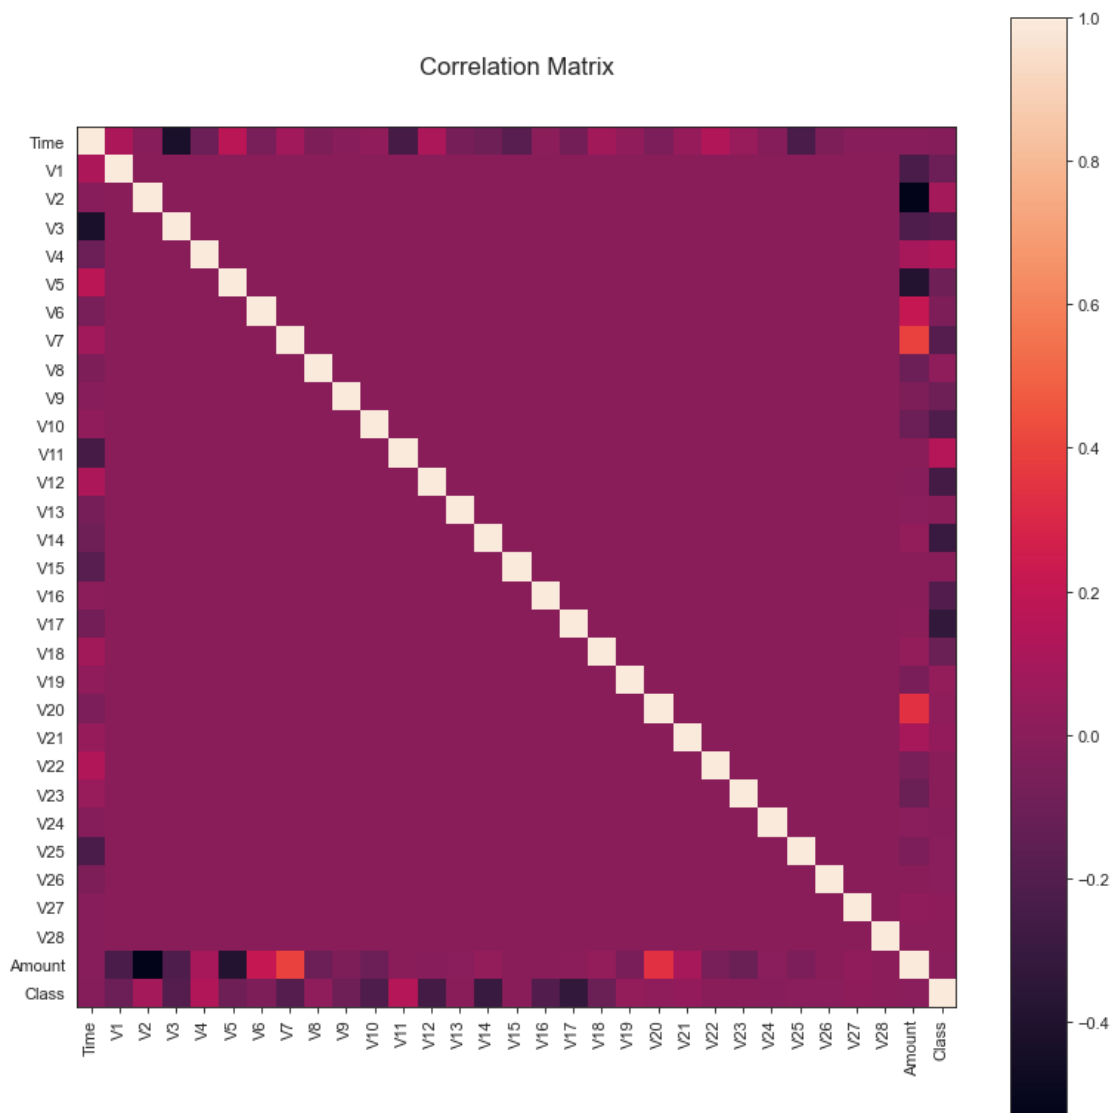
fig, axes = plt.subplots(2, 2, figsize=(25, 15))
fig.suptitle('Relations')

axes = axes.reshape(4,)

for i in range(len(max_corr_x)):
    sns.scatterplot(ax = axes[i],data = df, x= max_corr_x[i],y=max_corr_y[i])
```



```
[13]: #Correlation Matrix
df = df[[col for col in df if df[col].nunique() > 1]] # keep columns where
↳ there are more than 1 unique values
corr = df.corr()
plt.figure(num=None, figsize=(12, 12), dpi=80, facecolor='w', edgecolor='k')
corrMat = plt.matshow(corr, fignum = 1)
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
plt.gca().xaxis.tick_bottom()
plt.colorbar(corrMat)
plt.title(f'Correlation Matrix', fontsize=15)
plt.show()
```



```
[14]: df['Class'].value_counts()
```

```
[14]: 0    284315  
      1      492  
      Name: Class, dtype: int64
```

0.1.1 Train and Test

```
[15]: # Train and test data  
x=df.drop(columns=["Time", "Class"], axis="columns")  
y=df.Class
```

```
[16]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.  
      ↪25,random_state=42)
```

```
[17]: # Details of training dataset  
print("Shape of x_train dataset: ", x_train.shape)  
print("Shape of y_train dataset: ", y_train.shape)  
print("Shape of x_test dataset: ", x_test.shape)  
print("Shape of y_test dataset: ", y_test.shape)  
print("Before OverSampling, counts of label '1': {}".format(sum(y_train==1)))  
print("Before OverSampling, counts of label '0': {} \n".format(sum(y_train==0)))
```

```
Shape of x_train dataset: (213605, 29)  
Shape of y_train dataset: (213605,)  
Shape of x_test dataset: (71202, 29)  
Shape of y_test dataset: (71202,)  
Before OverSampling, counts of label '1': 379  
Before OverSampling, counts of label '0': 213226
```

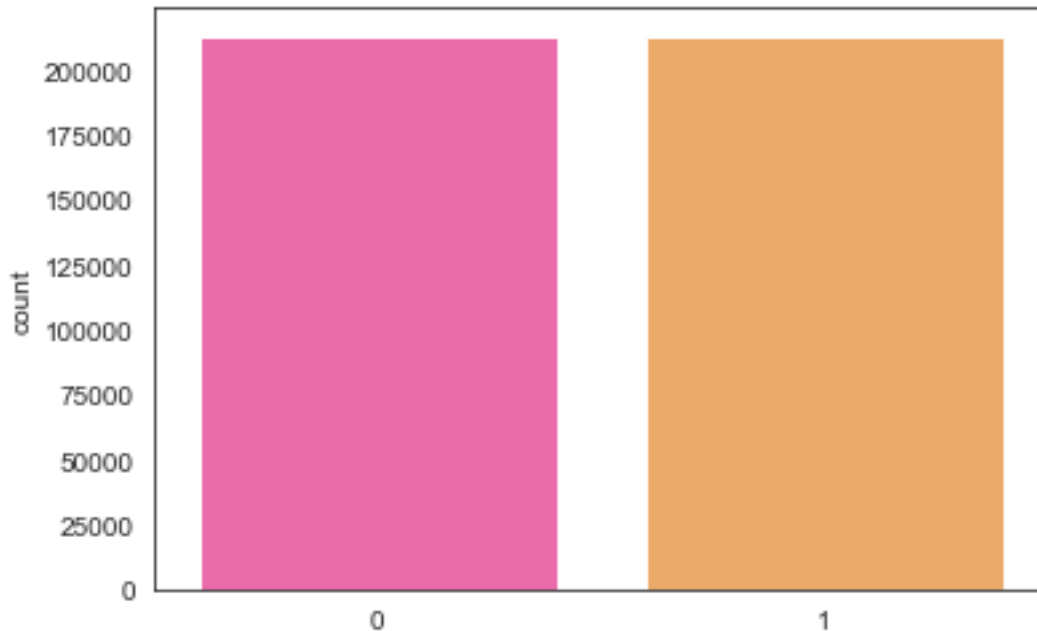
```
[18]: # Oversample the training dataset  
sm = SMOTE(random_state=2)  
x_train_s, y_train_s = sm.fit_resample(x_train, y_train.ravel())  
  
print('After OverSampling, the shape of train_x: {}'.format(x_train_s.shape))  
print('After OverSampling, the shape of train_y: {} \n'.format(y_train_s.shape))  
  
print("After OverSampling, counts of label '1', %: {}".format(sum(y_train_s==1)/  
      ↪len(y_train_s)*100.0,2))  
print("After OverSampling, counts of label '0', %: {}".format(sum(y_train_s==0)/  
      ↪len(y_train_s)*100.0,2))  
  
sns.countplot(x=y_train_s, data=df, palette='spring')
```

```
After OverSampling, the shape of train_x: (426452, 29)  
After OverSampling, the shape of train_y: (426452,)
```

After OverSampling, counts of label '1', %: 50.0

After OverSampling, counts of label '0', %: 50.0

[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8919687ca0>



[19]: *# Feature selection using Variance Threshold with threshold of 0.5*

```
var = VarianceThreshold(threshold=.5)
var.fit(x_train_s,y_train_s)
x_train_var=var.transform(x_train_s)
x_test_var=var.transform(x_test)
x_train_var.shape
```

[19]: (426452, 25)

[20]: *# Alternate way to perform feature selection and display the features*

```
def variance_threshold_selector(data, threshold=0.5):
    selector = VarianceThreshold(threshold)
    selector.fit(data)
    return data[data.columns[selector.get_support(indices=True)]]
variance_threshold_selector(x_train_s, 0.5)
```

[20]:

	V1	V2	V3	V4	V5	V6	V7	\
0	-1.648591	1.228130	1.370169	-1.735542	-0.029455	-0.484129	0.918645	
1	-0.234775	-0.493269	1.236728	-2.338793	-1.176733	0.885733	-1.960981	
2	1.134626	-0.774460	-0.163390	-0.533358	-0.604555	-0.244482	-0.212682	

```

3      0.069514  1.017753  1.033117  1.384376  0.223233 -0.310845  0.597287
4      -0.199441  0.610092 -0.114437  0.256565  2.290752  4.008475 -0.123530
...
426447  1.065197  0.639987  0.203695  3.005300 -0.066426  0.058743 -0.365901
426448 -6.689924  2.108646 -7.070325  5.133202 -2.313166 -2.003876 -8.549248
426449 -0.415163 -0.973361 -2.162616  2.541888  0.707195 -0.199787  1.793974
426450  0.172346  0.879474  1.678842  3.136607 -0.697444  1.213229 -1.449331
426451 -2.160564  1.424873 -0.282192  2.005992 -0.628338  0.148190 -1.810695

```

```

          V8      V9      V10  ...      V16      V17      V18  \
0      -0.438750  0.982144  1.241635  ...  0.664548 -1.280961  0.184568
1      -2.363412 -2.694774  0.360215  ... -0.163459  0.562423 -0.577032
2       0.040782 -1.136627  0.792009  ... -1.371503  0.020165  0.796223
3      -0.127658 -0.701533  0.070739  ... -0.507737 -0.024208  0.371960
4       1.038374 -0.075846  0.030453  ... -0.541172 -0.174950  0.355749
...
426447  0.318517 -0.139786 -0.311478  ...  1.110343  1.651054  1.223993
426448  0.674817 -3.590065 -8.093407  ... -7.755543 -12.704684 -4.443473
426449 -0.182592 -0.794665 -0.302585  ... -0.843149  1.686313  1.296810
426450 -1.552406 -0.848937  0.131740  ...  0.278056  0.626466  0.512756
426451  0.490887 -0.818383 -1.778571  ... -2.974153 -4.531439 -1.914134

```

```

          V19      V20      V21      V22      V23      V27      Amount
0      -0.331603  0.384201 -0.218076 -0.203458 -0.213015 -0.262968  38.420000
1      -1.635634  0.364679 -1.495358 -0.083066  0.074612  0.089293  61.200000
2      -0.519459 -0.396476 -0.684454 -1.855269  0.171997 -0.061178 110.950000
3       1.561447  0.148760  0.097023  0.369957 -0.219266  0.114440  10.000000
4       1.375281  0.292972 -0.019733  0.165463 -0.080978  0.481769  22.000000
...
426447 -0.872929 -0.193495 -0.087659 -0.084135 -0.126067  0.046251  1.725917
426448  2.498435  0.070287  0.067269  0.467306 -0.642386  1.339039  0.864189
426449  2.017400  1.171567  0.451141  0.535937  1.149598 -0.019759 453.619063
426450  0.001894  0.495637 -0.788119  0.929233 -0.117767  0.112793  0.339493
426451  0.606703  0.222745  0.425152  0.368008 -0.397128 -0.335962 17.827416

```

[426452 rows x 25 columns]

```
[21]: varth_features=var.get_support()
varth_features
```

```
[21]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
         True,  True,  True,  True,  True,  True,  True,  True,  True,
         True,  True,  True,  True,  True, False, False, False,  True,
         False,  True])
```

```
[22]: # Feature selection using SelectKBest feature selection
skbest = SelectKBest(k=10)
```

```

skbest.fit(x_train_s,y_train_s)
x_train_skbest=skbest.transform(x_train_s)
x_test_skbest=skbest.transform(x_test)
x_train_skbest.shape

```

[22]: (426452, 10)

```

[23]: kbest_features=skbest.get_support()
kbest_features

```

[23]: array([False, True, True, True, False, False, False, False, True,
 True, True, True, False, True, False, True, True, False,
 False, False, False, False, False, False, False, False, False,
 False, False])

```

[24]: # 10 best features using SelectKBest
best_features = SelectKBest(score_func=f_classif, k=10)
fit = best_features.fit(x_train_s,y_train_s)
df_scores = pd.DataFrame(fit.scores_)
df_columns = pd.DataFrame(x_train_s.columns)
feature_scores = pd.concat([df_columns, df_scores],axis=1)
feature_scores.columns = ['Feature_Name', 'Score'] # name output columns
print(feature_scores.nlargest(10, 'Score')) # print 10 best features

```

	Feature_Name	Score
13	V14	634408.764309
3	V4	477637.117795
10	V11	422096.275092
11	V12	415541.515783
9	V10	307081.858759
15	V16	240824.011003
8	V9	219463.452793
2	V3	203528.966849
16	V17	202096.481017
1	V2	151952.766954

```

[25]: # calculate precision recall area under curve
def preci_auc(y_true, pred_prob):
    # calculate precision-recall curve
    p, r, _ = precision_recall_curve(y_true, pred_prob)
    # calculate area under curve
    return auc(r, p)

```

```

[26]: # Evaluate a model
def evaluate_model(x, y, model):
    # Define evaluation procedure
    CV = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
    # Define the model evaluation the metric

```

```

metric = make_scorer(preci_auc, needs_proba=True)
# Evaluate model
scores = cross_val_score(model, x, y, scoring='roc_auc', cv=CV, n_jobs=-1)
return scores

```

```

[27]: # define the reference model
model = DummyClassifier(strategy='constant', constant=1)
# Evaluate the model
scores = evaluate_model(x_train_skbest, y_train_s, model)
# summarize performance
print('Mean area under curve: %.3f (%.3f)' % (mean(scores), std(scores)))

```

Mean area under curve: 0.500 (0.000)

```

[28]: # Normalize the input
scaler = StandardScaler()
scaler.fit(x_train_skbest)
x_train_norm = scaler.transform(x_train_skbest)
x_test_norm = scaler.transform(x_test_skbest)

```

```

[29]: def model_val(x, y, classifier, scor, show):
    x = np.array(x)
    y = np.array(y)

    scores = cross_val_score(classifier, x, y, scoring=scor)

    if show == True:
        print("Score: {:.2f} (+/- {:.2f})".format(scores.mean(), scores.std()))

    return scores.mean()

```

```

[30]: # List of models
rfc = RandomForestClassifier()
ctc = DecisionTreeClassifier()
sglc = SGDClassifier()
lr = LogisticRegression()

model = []
score = []

# Check model score
for classifier in (rfc, ctc, sglc, lr):
    model.append(classifier.__class__.__name__)
    score.append(model_val(x_train_norm, y_train_s, classifier, scor='roc_auc',
→ show=True))

pd.DataFrame(data=score, index=model, columns=['roc_auc'])

```



```
Score: 1.00 (+/- 0.00)
Score: 1.00 (+/- 0.00)
Score: 0.99 (+/- 0.00)
Score: 0.99 (+/- 0.00)
```

```
[30]:                                     roc_auc
RandomForestClassifier 0.999975
DecisionTreeClassifier 0.997339
SGDClassifier          0.990387
LogisticRegression     0.990506
```

0.1.2 Random Forest Model Evaluation

```
[31]: pipeline_rf = Pipeline([
      ('model', RandomForestClassifier(n_jobs=-1, random_state=1))
    ])
parm_gridscv_rf = {'model__n_estimators': [75]}
grid_rf = GridSearchCV(estimator=pipeline_rf, param_grid=parm_gridscv_rf,
    ↪scoring='roc_auc', n_jobs=-1,
                        pre_dispatch='2*n_jobs', cv=5, verbose=1,
    ↪return_train_score=False)
grid_rf.fit(x_train_norm, y_train_s)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[31]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('model',
                                             RandomForestClassifier(n_jobs=-1,
                                                                     random_state=1))])),
      n_jobs=-1, param_grid={'model__n_estimators': [75]},
      scoring='roc_auc', verbose=1)
```

```
[32]: pd.DataFrame(grid_rf.cv_results_)
```

```
[32]:   mean_fit_time  std_fit_time  mean_score_time  std_score_time \
0      122.154728    0.336013         0.27644         0.058226

   param_model__n_estimators      params  split0_test_score \
0                        75  {'model__n_estimators': 75}      0.999974

   split1_test_score  split2_test_score  split3_test_score  split4_test_score \
0           0.999987           0.999925           0.999961           0.999959

   mean_test_score  std_test_score  rank_test_score
0           0.999961           0.000021              1
```

```
[33]: grid_rf.best_score_, grid_rf.best_params_
```

```
[33]: (0.9999611455402032, {'model__n_estimators': 75})
```

0.1.3 Test Random Forest model

```
[34]: # Prediction for the test set
y_pred_test = grid_rf.predict(x_test_norm)
# Decimal places based on number of samples
dec = np.int64(np.ceil(np.log10(len(y_test))))

print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred_test), '\n')

print('Classification report')
print(classification_report(y_test, y_pred_test, digits=dec))

print('Scalar Metrics')
format_str = '%13s = %%.%if' % dec
if y_test.nunique() <= 2: # metrics for binary classification
    try:
        y_score = grid_rf.predict_proba(x_test_norm)[:,-1]
    except:
        y_score = grid_rf.decision_function(x_test_norm)
    print(format_str % ('AUROC', roc_auc_score(y_test, y_score)))
```

Confusion Matrix

```
[[71049   40]
 [   13  100]]
```


Classification report

	precision	recall	f1-score	support
0	0.99982	0.99944	0.99963	71089
1	0.71429	0.88496	0.79051	113
accuracy			0.99926	71202
macro avg	0.85705	0.94220	0.89507	71202
weighted avg	0.99936	0.99926	0.99930	71202

Scalar Metrics

AUROC = 0.98299

```
[35]: # Plot confusion matrix
con_mat=confusion_matrix(y_test,y_pred_test,labels=[0,1])
cmatrix=ConfusionMatrixDisplay(confusion_matrix=con_mat,display_labels=["Not_
↪Fraud", "Fraud"])
cmatrix.plot()
plt.title("Confusion Matrix")
plt.show()
```



```
[36]: log_loss(y_test, y_pred_test)
```

```
[36]: 0.025709771254003325
```

0.1.4 Logistic Regression Model Evaluation

```
[37]: # Logistic regression model with different C values
parameters = {
    'tol': [0.00001, 0.0001, 0.001],
    'C': [1, 50, 100]
}

lgr = GridSearchCV(LogisticRegression(random_state=101, n_jobs=1,
    ↪max_iter=1000),
                    param_grid=parameters,
                    cv=3,
                    n_jobs=1,
                    scoring='roc_auc'
                )
lgr.fit(x_train_norm, y_train_s)
clf = lgr.best_estimator_

print(lgr.best_estimator_)
print("The best classifier score:", lgr.best_score_)
```

LogisticRegression(C=1, max_iter=1000, n_jobs=1, random_state=101, tol=1e-05)
The best classifier score: 0.9905075847918005

0.1.5 Test Logistic Regression Model

```
[38]: y_pred_test1 = clf.predict(x_test_norm)
# Decimal places based on number of samples
dec = np.int64(np.ceil(np.log10(len(y_test))))

print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred_test1), '\n')

print('Classification report')
print(classification_report(y_test, y_pred_test1, digits=dec))

print('Scalar Metrics')
format_str = '%13s = %%.%if' % dec
if y_test.nunique() <= 2: # metrics for binary classification
    try:
        y_score1 = clf.predict_proba(x_test_norm)[: ,1]
    except:
        y_score1 = clf.decision_function(X_test_norm)
    print(format_str % ('AUROC', roc_auc_score(y_test, y_score1)))
```

Confusion Matrix

```
[[69651 1438]
 [    9  104]]
```


Classification report

	precision	recall	f1-score	support
0	0.99987	0.97977	0.98972	71089
1	0.06744	0.92035	0.12568	113
accuracy			0.97968	71202
macro avg	0.53366	0.95006	0.55770	71202
weighted avg	0.99839	0.97968	0.98835	71202

Scalar Metrics

AUROC = 0.97698

```
[39]: # Plot confusion matrix
con_mat1=confusion_matrix(y_test,y_pred_test1,labels=[0,1])
cmatrix1=ConfusionMatrixDisplay(confusion_matrix=con_mat1,display_labels=["Not_Fraud", "Fraud"])
cmatrix1.plot()
plt.title("Confusion Matrix")
plt.show()
```



```
[40]: log_loss(y_test, y_pred_test1)
```

```
[40]: 0.7019291489640804
```

0.1.6 Train and Test after balancing

```
[138]: df = df.sample(frac=1)

# Taking only few sample cases where length of sample is equal to number of
↳ fraud cases.
fraud_df = df.loc[df['Class'] == 1]
non_fraud_df = df.loc[df['Class'] == 0].sample(len(fraud_df))

normal_distributed_df = pd.concat([fraud_df, non_fraud_df])

# Shuffle dataframe rows
new_df = normal_distributed_df.sample(frac=1, random_state=42)

new_df.head()
```

```
[138]:
```

	Time	V1	V2	V3	V4	V5	V6	\
180249	124450.0	2.109715	-0.458784	-1.436015	-1.452778	-0.364377	-1.586309	
154697	102625.0	-4.221221	2.871121	-5.888716	6.890952	-3.404894	-1.154394	
35009	37917.0	1.384650	-1.376614	0.417809	-1.567336	-1.471850	0.004499	
230076	146179.0	-0.067672	4.251181	-6.540388	7.283657	0.513541	-2.635066	
197586	132086.0	-0.361428	1.133472	-2.971360	-0.283073	0.371452	-0.574680	


	V7	V8	V9	...	V21	V22	V23	\
180249	0.121582	-0.366016	1.753300	...	-0.260543	-0.541017	0.275271	
154697	-7.739928	2.851363	-2.507569	...	1.620591	1.567947	-0.578007	
35009	-1.299660	0.053219	-1.772205	...	-0.025284	0.148147	-0.171506	
230076	-1.865911	0.780272	-3.868248	...	0.415437	-0.469938	0.007128	
197586	4.031513	-0.934398	-0.768255	...	0.110815	0.563861	-0.408436	

	V24	V25	V26	V27	V28	Amount	Class
180249	-0.098500	-0.117366	-0.484246	-0.014693	-0.058168	4.69	0
154697	-0.059045	-1.829169	-0.072429	0.136734	-0.599848	7.59	1
35009	-0.507789	0.388308	-0.075662	0.032970	0.019813	79.00	0
230076	-0.388147	-0.493398	0.466468	0.566370	0.262990	0.77	1
197586	-0.880079	1.408392	-0.137402	-0.001250	-0.182751	480.72	1

[5 rows x 31 columns]

```
[139]: #Correlation Matrix
new_df = new_df[[col for col in new_df if new_df[col].nunique() > 1]] # keep
↳ columns where there are more than 1 unique values
corr = new_df.corr()
plt.figure(num=None, figsize=(12, 12), dpi=80, facecolor='w', edgecolor='k')
corrMat = plt.matshow(corr, fignum = 1)
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
```

```
plt.gca().xaxis.tick_bottom()
plt.colorbar(corrMat)
plt.title(f'Correlation Matrix', fontsize=15)
plt.show()
```



```
[140]: # Train and test data
x=new_df.drop(columns=["Time","Class"],axis="columns")
y=new_df.Class
```

```
[141]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.
    ↪25,random_state=42)
```

```
[142]: # Details of training dataset
print("Shape of x_train dataset: ", x_train.shape)
print("Shape of y_train dataset: ", y_train.shape)
print("Shape of x_test dataset: ", x_test.shape)
print("Shape of y_test dataset: ", y_test.shape)
print("Before OverSampling, counts of label '1': {}".format(sum(y_train==1)))
print("Before OverSampling, counts of label '0': {} \n".format(sum(y_train==0)))
```

```
Shape of x_train dataset: (738, 29)
Shape of y_train dataset: (738,)
Shape of x_test dataset: (246, 29)
Shape of y_test dataset: (246,)
Before OverSampling, counts of label '1': 355
Before OverSampling, counts of label '0': 383
```

```
[143]: # Oversample the training dataset
sm = SMOTE(random_state=2)
x_train_s, y_train_s = sm.fit_resample(x_train, y_train.ravel())

print('After OverSampling, the shape of train_x: {}'.format(x_train_s.shape))
print('After OverSampling, the shape of train_y: {} \n'.format(y_train_s.shape))


print("After OverSampling, counts of label '1', %: {}".format(sum(y_train_s==1)/
    ↳len(y_train_s)*100.0,2))
print("After OverSampling, counts of label '0', %: {}".format(sum(y_train_s==0)/
    ↳len(y_train_s)*100.0,2))

sns.countplot(x=y_train_s, data=new_df, palette='spring')
```

```
After OverSampling, the shape of train_x: (766, 29)
After OverSampling, the shape of train_y: (766,)
```

```
After OverSampling, counts of label '1', %: 50.0
After OverSampling, counts of label '0', %: 50.0
```

```
[143]: <matplotlib.axes._subplots.AxesSubplot at 0x7f88e9793220>
```

[144]: *# Feature selection using Variance Threshold with threshold of 0.5*

```
var = VarianceThreshold(threshold=.5)
var.fit(x_train_s,y_train_s)
x_train_var=var.transform(x_train_s)
x_test_var=var.transform(x_test)
x_train_var.shape
```

[144]: (766, 25)

[145]: *# Alternate way to perform feature selection and display the features*

```
def variance_threshold_selector(data, threshold=0.5):
    selector = VarianceThreshold(threshold)
    selector.fit(data)
    return data[data.columns[selector.get_support(indices=True)]]
variance_threshold_selector(x_train_s, 0.5)
```

[145]:

	V1	V2	V3	V4	V5	V6	V7 \
0	-1.396204	2.618584	-6.036770	3.552454	1.030091	-2.950358	-1.528506
1	-0.518521	1.629458	1.998444	2.897556	0.275262	0.357826	0.409205
2	-11.205461	7.914633	-13.987752	4.333341	-8.484970	-3.506561	-8.935243
3	-3.821939	5.667247	-9.244963	8.246147	-4.368286	-3.450735	-8.427378
4	-6.618211	3.835943	-6.316453	1.844111	-2.476892	-1.886718	-3.817495
..
761	-0.794559	5.026167	-8.053515	7.465753	0.114331	-2.344652	-3.116944
762	1.113945	2.989537	-4.925187	4.266583	2.868245	-1.387594	0.927837

```

763 -5.101215  8.938518 -15.821726  10.301614 -4.601440 -3.373624 -11.049532
764 -2.801712 -0.263848 -5.731595   2.295171 -1.876356  0.590591 -1.160552
765 -3.878726  2.531532 -2.946649   4.093946 -1.516382 -0.855728 -4.386300

```

```

      V8      V9      V10 ...      V16      V17      V18 \
0   0.189319 -1.433554 -5.569142 ... -2.497341 -1.588336  0.120289
1   0.044234 -0.102798  0.540645 ... -0.314980  0.979470 -0.289904
2   7.704449 -2.336584 -5.927359 ... -3.847293 -6.700637 -2.492616
3   2.305609 -5.338079 -12.011161 ... -12.105602 -21.338195 -8.045436
4   0.613470 -1.482121 -4.868747 ... -3.939384 -7.164430 -2.434672
..      ...      ...      ...      ...      ...      ...
761  1.487198 -4.550283 -5.570927 ... -2.391081 -2.354116  0.327073
762 -1.555537 -1.560760 -1.588009 ...  2.931121  5.723689  3.748270
763  5.368419 -5.673940 -11.652628 ... -10.152260 -13.687070 -4.993079
764  1.445675 -1.869726 -6.197320 ... -4.095083 -5.457750 -1.690061
765 -0.063074 -1.199092 -5.433314 ... -6.866600 -9.750700 -4.469437

```

```

      V19      V20      V21      V22      V23      V27      Amount
0   0.170144  0.031795  0.143177 -0.390176  0.356029  0.062655   1.000000
1   0.293955  0.179944 -0.300826 -0.550231  0.043216  0.235407   6.480000
2   0.469554  0.860912  0.942593 -0.987848 -0.279446  1.084023  99.990000
3   0.156015  1.115247  1.990520  0.083353 -0.062264  1.869570  75.860000
4   0.235227 -0.953827  1.636622  0.038727  0.278218 -2.042403  57.730000
..      ...      ...      ...      ...      ...      ...
761  0.793628  0.825869  0.579042 -0.469174 -0.005520  0.657755   0.770000
762 -1.793996 -0.356969  1.099571 -0.795293  0.292327  0.146284   0.807329
763  1.137330  1.449403  2.000051  0.200852  0.602423  1.591059   1.000000
764 -0.267847  2.196937  1.401291  0.840844  1.323173  0.696839  723.210000
765  0.434441  0.019369  1.467101 -0.073473 -0.036148 -0.112660   0.929997

```

[766 rows x 25 columns]

```
[146]: varth_features=var.get_support()
varth_features
```

```
[146]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
         True,  True,  True,  True,  True,  True,  True,  True,  True,
         True,  True,  True,  True,  True, False, False, False,  True,
        False,  True])
```

```
[147]: # Feature selection using SelectKBest feature selection
skbest = SelectKBest(k=10)
skbest.fit(x_train_s,y_train_s)
x_train_skbest=skbest.transform(x_train_s)
x_test_skbest=skbest.transform(x_test)
x_train_skbest.shape
```

[147]: (766, 10)

```
[148]: kbest_features=skbest.get_support()
kbest_features
```

```
[148]: array([False,  True,  True,  True, False, False, False, False,  True,
         True,  True,  True, False,  True, False,  True,  True, False,
        False, False, False, False, False, False, False, False, False,
        False, False])
```

```
[149]: # 10 best features using SelectKBest
best_features = SelectKBest(score_func=f_classif, k=10)
fit = best_features.fit(x_train_s,y_train_s)
df_scores = pd.DataFrame(fit.scores_)
df_columns = pd.DataFrame(x_train_s.columns)
feature_scores = pd.concat([df_columns, df_scores],axis=1)
feature_scores.columns = ['Feature_Name', 'Score'] # name output columns
print(feature_scores.nlargest(10, 'Score')) # print 10 best features
```

	Feature_Name	Score
13	V14	961.878974
3	V4	768.888297
10	V11	703.241720
11	V12	661.244182
9	V10	534.163705
15	V16	453.489625
2	V3	390.362289
16	V17	384.380278
8	V9	369.572827
1	V2	247.774452

```
[150]: # calculate precision recall area under curve
def preci_auc(y_true, pred_prob):
    # calculate precision-recall curve
    p, r, _ = precision_recall_curve(y_true, pred_prob)
    # calculate area under curve
    return auc(r, p)
```

```
[151]: # Evaluate a model
def evaluate_model(x, y, model):
    # Define evaluation procedure
    CV = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
    # Define the model evaluation the metric
    metric = make_scorer(preci_auc, needs_proba=True)
    # Evaluate model
    scores = cross_val_score(model, x, y, scoring='roc_auc', cv=CV, n_jobs=-1)
    return scores
```

```
[152]: # define the reference model
model = DummyClassifier(strategy='constant', constant=1)
# Evaluate the model
scores = evaluate_model(x_train_skbtest, y_train_s, model)
# summarize performance
print('Mean area under curve: %.3f (%.3f)' % (mean(scores), std(scores)))
```

Mean area under curve: 0.500 (0.000)

```
[153]: # Normalize the input
scaler = StandardScaler()
scaler.fit(x_train_skbtest)
x_train_norm = scaler.transform(x_train_skbtest)
x_test_norm = scaler.transform(x_test_skbtest)
```

```
[154]: def model_val(x, y, classifier, scor, show):
    x = np.array(x)
    y = np.array(y)

    scores = cross_val_score(classifier, x, y, scoring=scor)

    if show == True:
        print("Score: {:.2f} (+/- {:.2f})".format(scores.mean(), scores.std()))

    return scores.mean()
```

```
[155]: # List of models
rfc = RandomForestClassifier()
ctc = DecisionTreeClassifier()
sglc = SGDClassifier()
lr = LogisticRegression()

model = []
score = []

# Check model score
for classifier in (rfc, ctc, sglc, lr):
    model.append(classifier.__class__.__name__)
    score.append(model_val(x_train_norm, y_train_s, classifier, scor='roc_auc',
↪ show=True))

pd.DataFrame(data=score, index=model, columns=['roc_auc'])
```

Score: 0.97 (+/- 0.01)

Score: 0.89 (+/- 0.02)

Score: 0.98 (+/- 0.01)

Score: 0.98 (+/- 0.01)

```
[155]:
```

	roc_auc
RandomForestClassifier	0.971959
DecisionTreeClassifier	0.889183
SGDClassifier	0.975939
LogisticRegression	0.978982

0.1.7 Random Forest Model Evaluation after balancing

```
[156]: pipeline_rf = Pipeline([
    ('model', RandomForestClassifier(n_jobs=-1, random_state=1))
])
param_gridscv_rf = {'model__n_estimators': [75]}
grid_rf = GridSearchCV(estimator=pipeline_rf, param_grid=param_gridscv_rf,
    ↪scoring='roc_auc', n_jobs=-1,
    pre_dispatch='2*n_jobs', cv=5, verbose=1,
    ↪return_train_score=False)
grid_rf.fit(x_train_norm, y_train_s)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[156]: GridSearchCV(cv=5,
    estimator=Pipeline(steps=[('model',
                                RandomForestClassifier(n_jobs=-1,
random_state=1))])),
    n_jobs=-1, param_grid={'model__n_estimators': [75]},
    scoring='roc_auc', verbose=1)
```

```
[157]: pd.DataFrame(grid_rf.cv_results_)
```

```
[157]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	0.2282	0.016534	0.022987	0.001361	

	param_model__n_estimators	params	split0_test_score	\
0	75	{'model__n_estimators': 75}	0.980098	

	split1_test_score	split2_test_score	split3_test_score	split4_test_score	\
0	0.965311	0.969754	0.983424	0.968301	

	mean_test_score	std_test_score	rank_test_score
0	0.973378	0.007072	1

```
[158]: grid_rf.best_score_, grid_rf.best_params_
```

```
[158]: (0.9733777329983756, {'model__n_estimators': 75})
```

0.1.8 Test Random Forest model after balancing

```
[159]: # Prediction for the test set
y_pred_test = grid_rf.predict(x_test_norm)
# Decimal places based on number of samples
dec = np.int64(np.ceil(np.log10(len(y_test))))

print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred_test), '\n')

print('Classification report')
print(classification_report(y_test, y_pred_test, digits=dec))

print('Scalar Metrics')
format_str = '%13s = %%.%if' % dec
if y_test.nunique() <= 2: # metrics for binary classification
    try:
        y_score = grid_rf.predict_proba(x_test_norm)[:,-1]
    except:
        y_score = grid_rf.decision_function(x_test_norm)
    print(format_str % ('AUROC', roc_auc_score(y_test, y_score)))
```

Confusion Matrix

```
[[104   5]
 [ 11 126]]
```


Classification report

	precision	recall	f1-score	support
0	0.904	0.954	0.929	109
1	0.962	0.920	0.940	137
accuracy			0.935	246
macro avg	0.933	0.937	0.934	246
weighted avg	0.936	0.935	0.935	246

Scalar Metrics

AUROC = 0.973

```
[160]: # Plot confusion matrix
con_mat=confusion_matrix(y_test,y_pred_test,labels=[0,1])
cmatrix=ConfusionMatrixDisplay(confusion_matrix=con_mat,display_labels=["Not_
↪Fraud", "Fraud"])
cmatrix.plot()
plt.title("Confusion Matrix")
plt.show()
```



```
[161]: log_loss(y_test, y_pred_test)
```

```
[161]: 2.2464407329500897
```

0.1.9 Logistic Regression Model Evaluation after balancing

```
[162]: # Logistic regression model with different C values
parameters = {
    'tol': [0.00001, 0.0001, 0.001],
    'C': [1, 50, 100]
}

lgr = GridSearchCV(LogisticRegression(random_state=101, n_jobs=1,
    ↪max_iter=1000),
                    param_grid=parameters,
                    cv=3,
                    n_jobs=1,
                    scoring='roc_auc'
                    )
lgr.fit(x_train_norm, y_train_s)
clf = lgr.best_estimator_

print(lgr.best_estimator_)
print("The best classifier score:", lgr.best_score_)
```

LogisticRegression(C=1, max_iter=1000, n_jobs=1, random_state=101, tol=1e-05)
The best classifier score: 0.9782445712352362

0.1.10 Test Logistic Regression Model after balancing

```
[163]: y_pred_test1 = clf.predict(x_test_norm)
# Decimal places based on number of samples
dec = np.int64(np.ceil(np.log10(len(y_test))))

print('Confusion Matrix')
print(confusion_matrix(y_test, y_pred_test1), '\n')

print('Classification report')
print(classification_report(y_test, y_pred_test1, digits=dec))

print('Scalar Metrics')
format_str = '%13s = %%.%if' % dec
if y_test.nunique() <= 2: # metrics for binary classification
    try:
        y_score1 = clf.predict_proba(x_test_norm)[:,-1]
    except:
        y_score1 = clf.decision_function(X_test_norm)
    print(format_str % ('AUROC', roc_auc_score(y_test, y_score1)))
```

Confusion Matrix

```
[[106   3]
 [ 13 124]]
```


Classification report

	precision	recall	f1-score	support
0	0.891	0.972	0.930	109
1	0.976	0.905	0.939	137
accuracy			0.935	246
macro avg	0.934	0.939	0.935	246
weighted avg	0.938	0.935	0.935	246

Scalar Metrics

AUROC = 0.985

```
[164]: # Plot confusion matrix
con_mat1=confusion_matrix(y_test,y_pred_test1,labels=[0,1])
cmatrix1=ConfusionMatrixDisplay(confusion_matrix=con_mat1,display_labels=["Not_Fraud", "Fraud"])
cmatrix1.plot()
plt.title("Confusion Matrix")
plt.show()
```



```
[165]: log_loss(y_test, y_pred_test1)
```

```
[165]: 2.246434232157974
```

[]: